

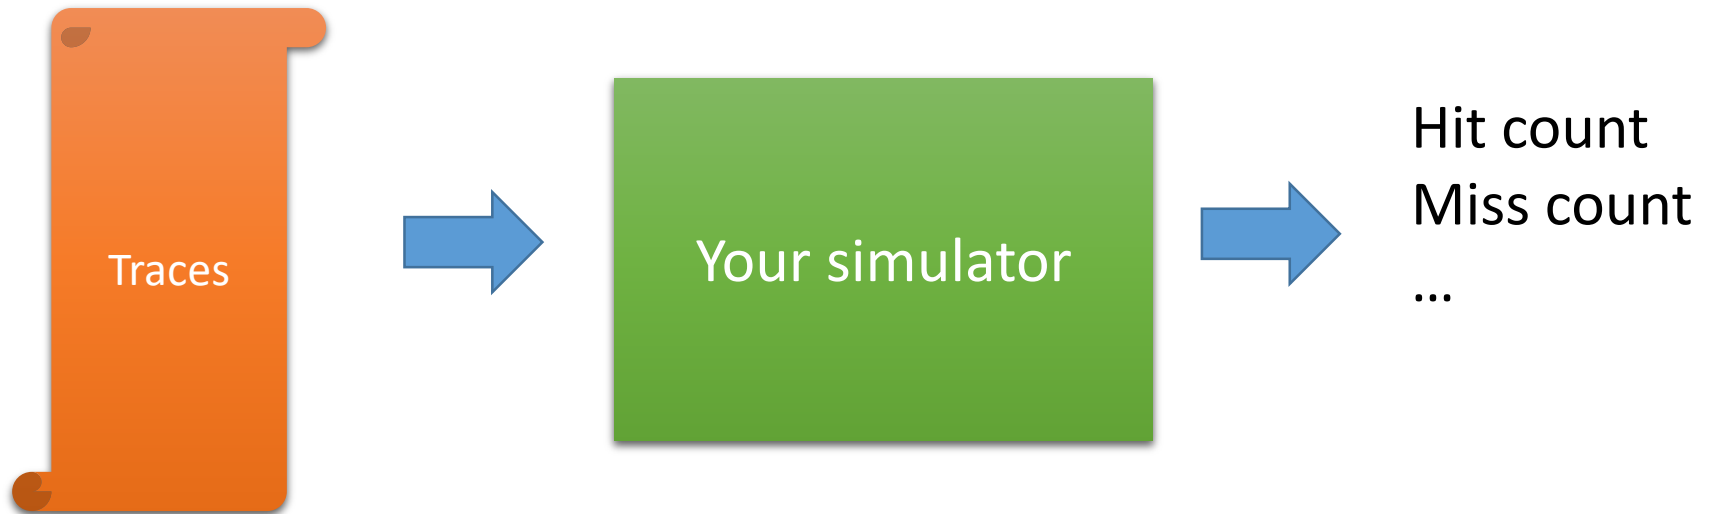
# Operating Systems

## Programming Assignment #5

Page Replacement Simulation: LRU and LFU

Prof. Li-Pin Chang, NCTU

# Simulation



# Trace File Format

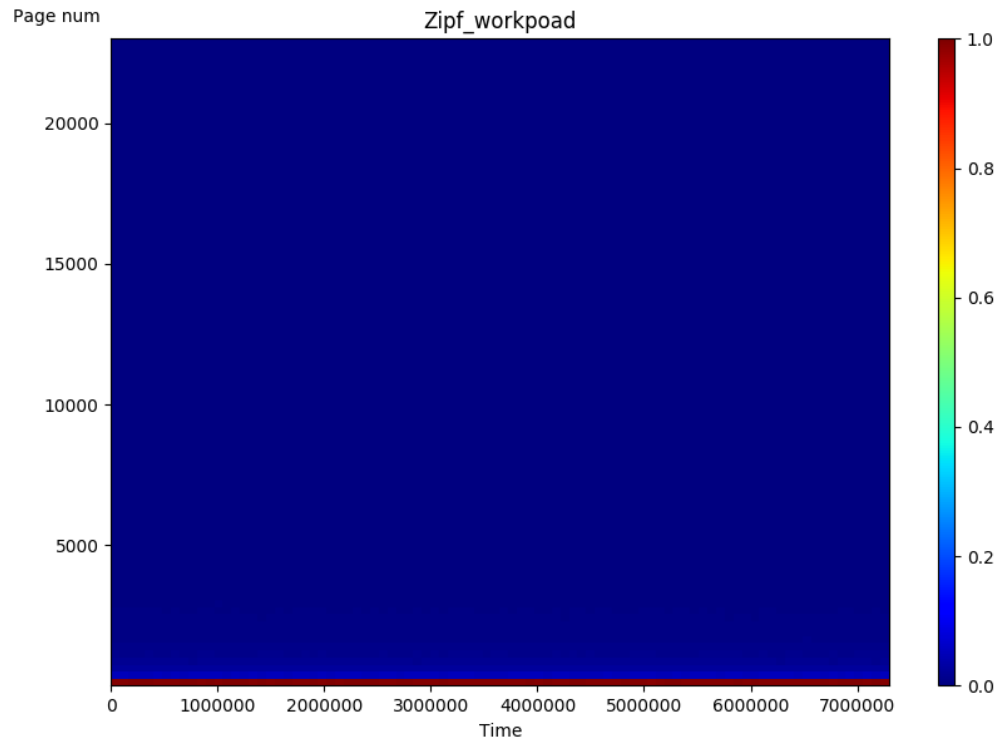
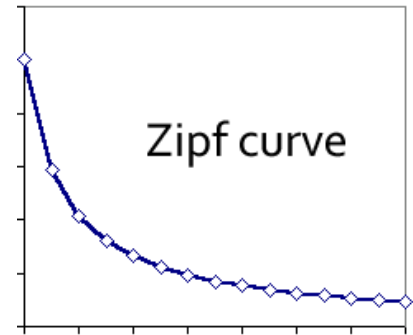
- Page number of referenced pages

```
1003  
1003  
9340  
1243  
1108  
1786  
1066  
1312  
1000  
1000  
1213  
1249  
2116
```

Unsigned integer

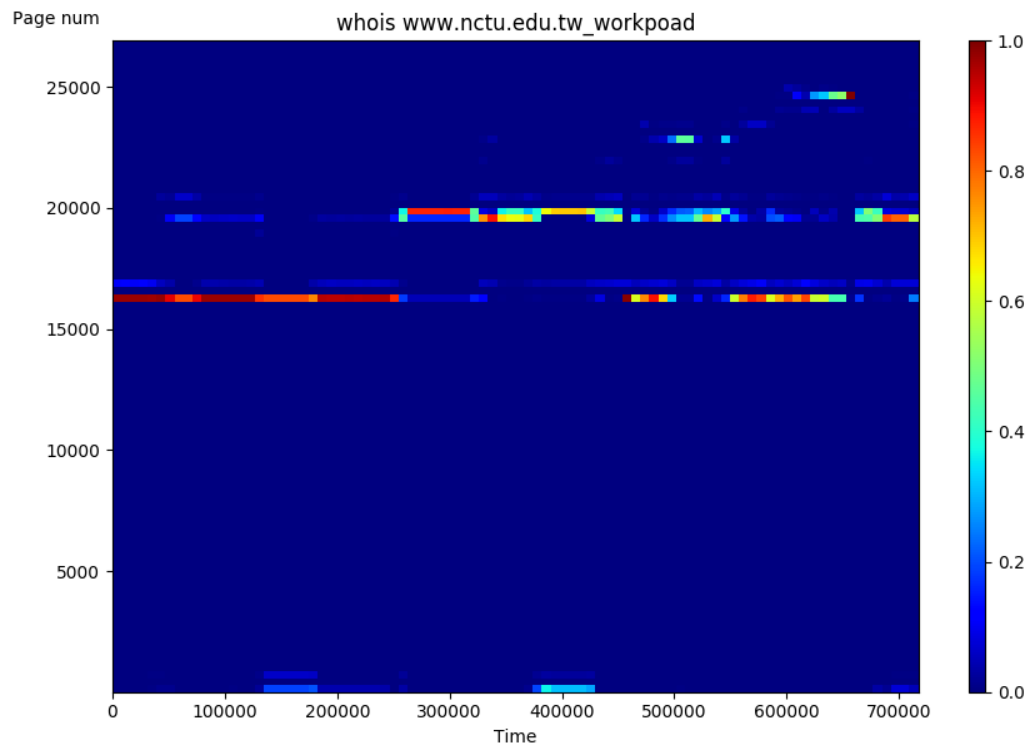
# Test Workload (I)

- Zipfian distribution



# Test Workload (II)

- “whois www.nctu.edu.tw”



# Page Replacement(LFU)

- Example: Frame #=2

4001 (miss)

LFU	
4001	

Ref count		
4000	4001	4002
0	1	0

4001 (hit)

4001	
------	--

0	2	0
---	---	---

4000 (miss)

4000	4001
------	------

1	2	0
---	---	---

4002 (miss)

4000 ←	4002	4001
--------	------	------

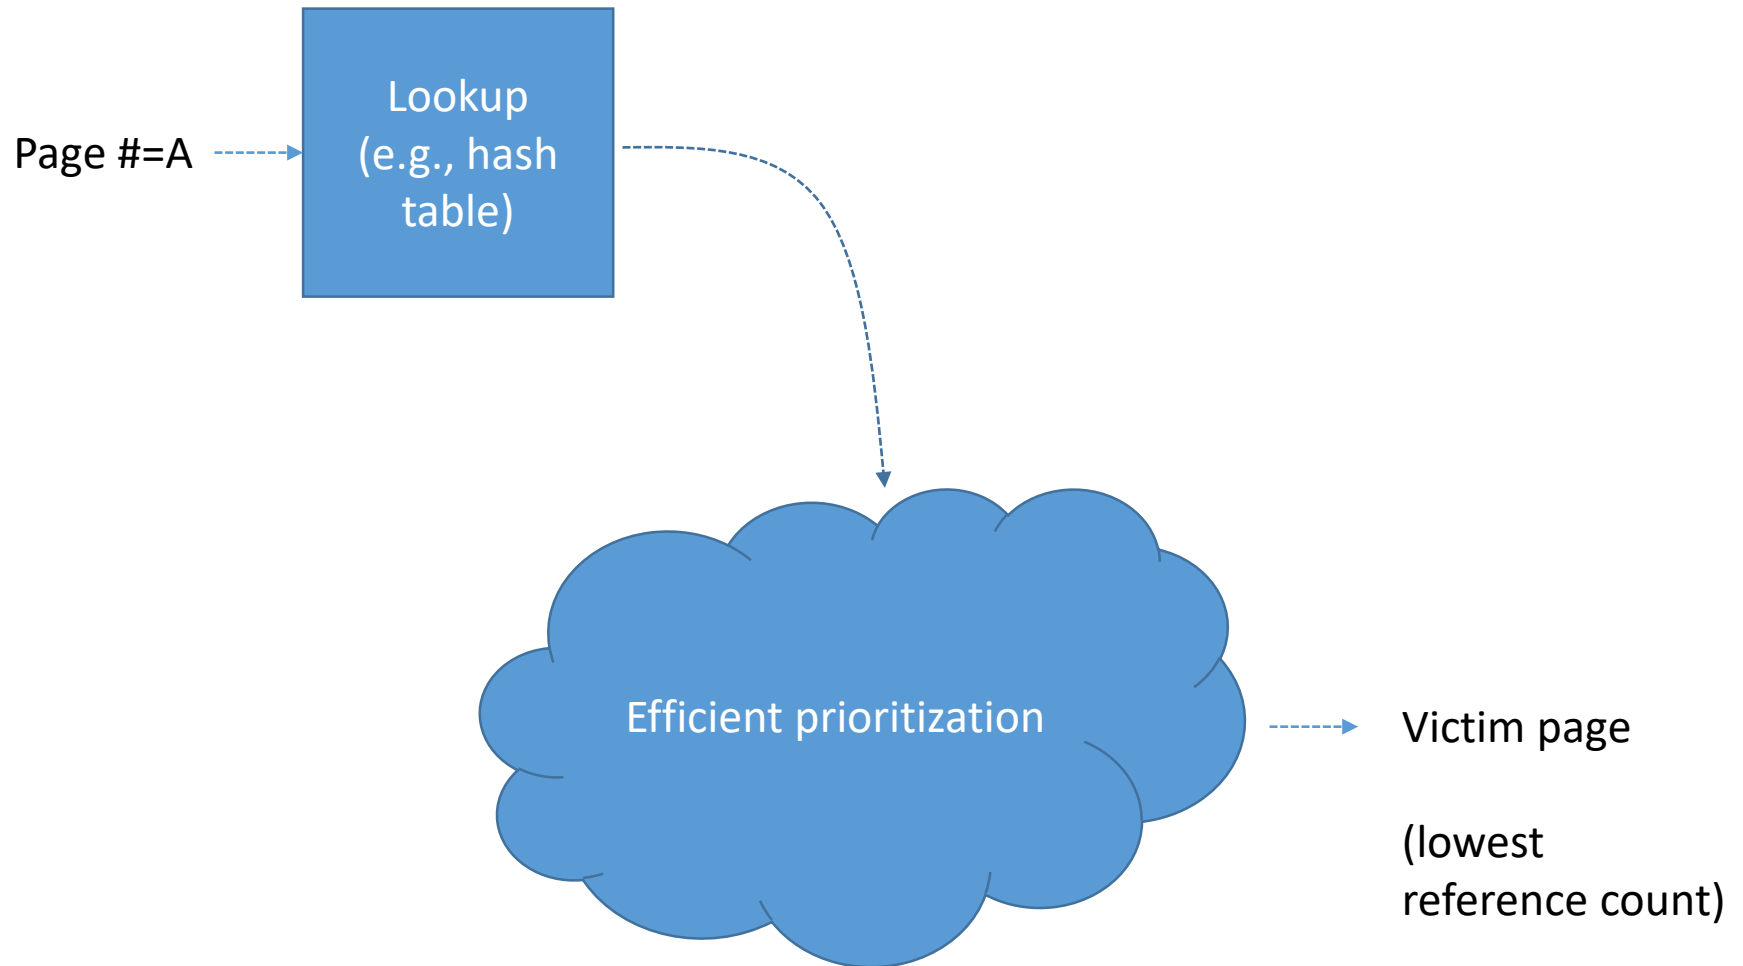
0	2	1
---	---	---

4000 (miss)

4002 ←	4000	4001
--------	------	------

1	2	0
---	---	---

# Simulator Structure (LFU)



# Page Replacement(LRU)

- Example: Frame #=2

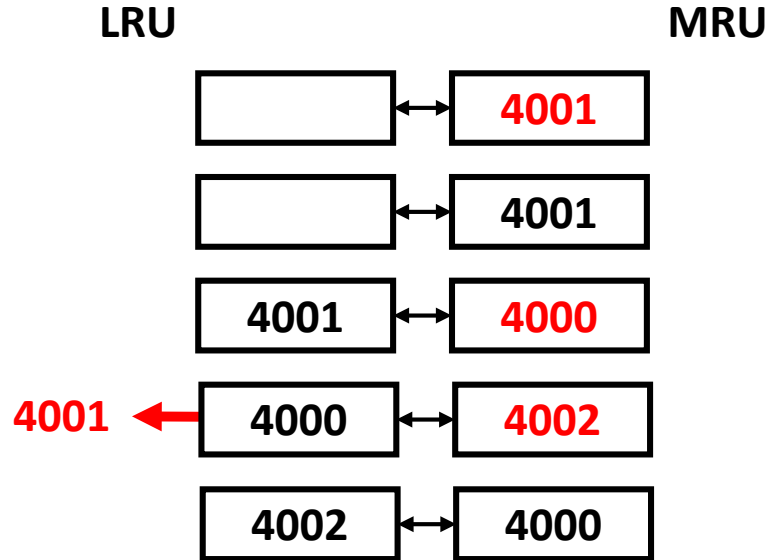
4001 (miss)

4001 (hit)

4000 (miss)

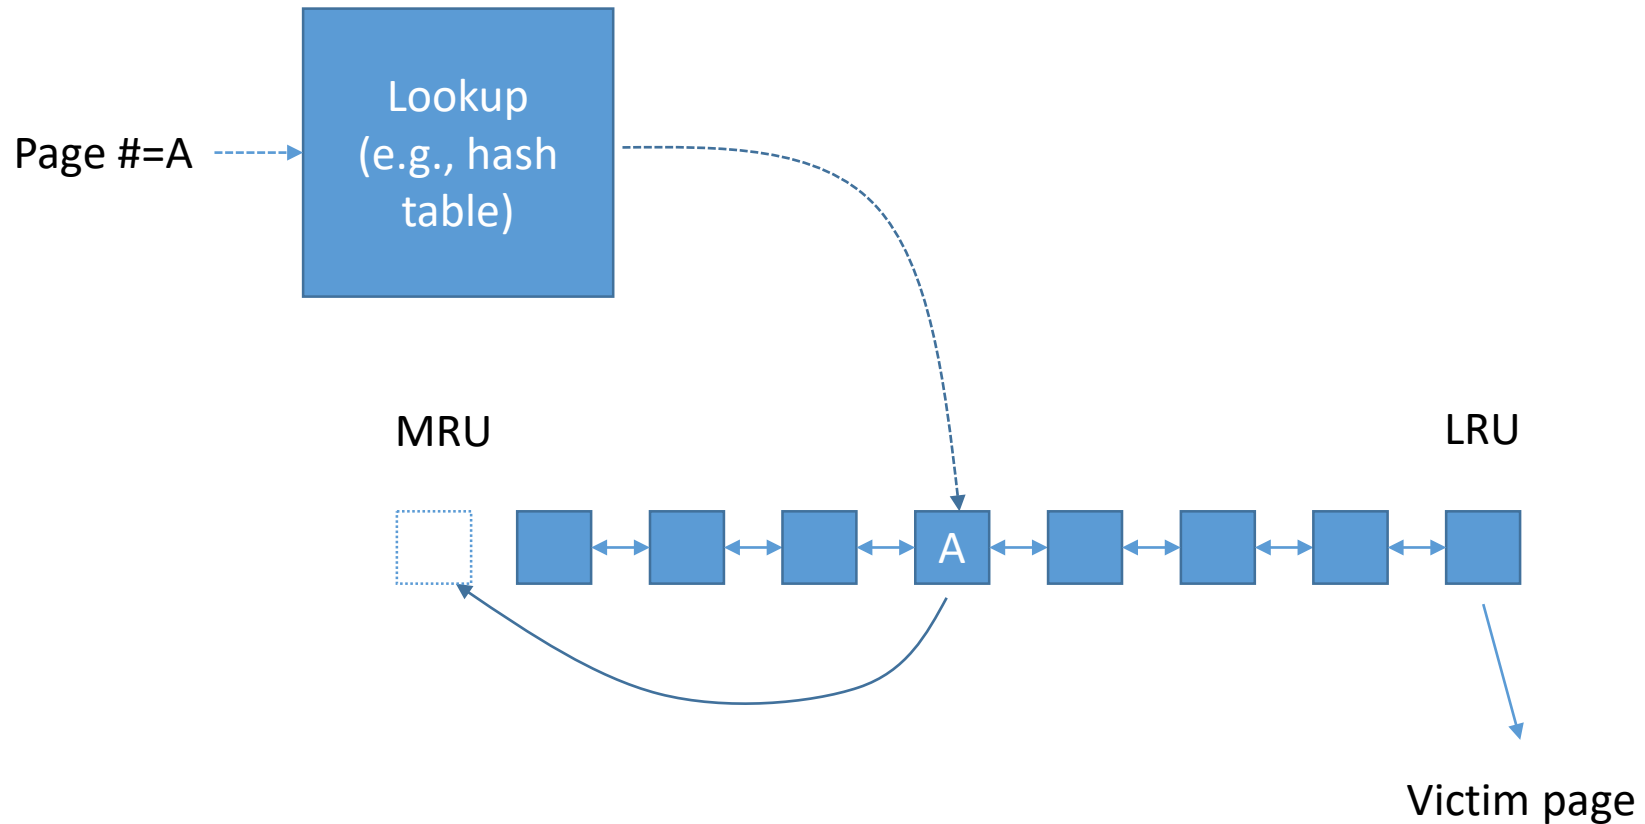
4002 (miss)

4000 (hit)





# Simulator Structure (LRU)



# Page Cache Operations

- Page lookup
  - Check whether or a new reference is a hit or a miss
  - Hash tables, binary search trees, skip lists....
- Do not use linear search!!!
  - You will receive a grade penalty if you do
  - Implement your own search, or reuse any existing libraries/classes for searching
  - TAs will examine your code
  - Duplication in this part does not count

# Victim selection

- LFU
  - The least frequently used page
  - If two pages have the same access count, the page having a smaller reference sequence number is replaced
  - You may need to store the reference sequence number when a page is added to the page cache
- LRU
  - The least recently used page

# Procedure

1. Algorithm=LFU
2. For (frame # = 64, 128, 256, and 512) do
  - Read the trace file
  - Run simulation
  - Print out the hit count, miss count, page fault ratio
3. Print out the total elapsed time of Step 2
4. Algorithm=LRU
5. For (frame # = 64, 128, 256, and 512) do
  - Read the trace file
  - Run simulation
  - Print out the hit count, miss count, page fault ratio
6. Print out the total elapsed time of Step 5

# The scenario of program

## Output format

LFU policy: (\n)

Frame	(\t)	Hit	(\t\t)	Miss	(\t\t)	Page fault ratio(\n)
64	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f(\n)
128	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f(\n)
256	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f(\n)
512	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f(\n)

Total elapsed time %.4f sec(\n) (\n)

LRU policy: (\n)

Frame	(\t)	Hit	(\t\t)	Miss	(\t\t)	Page fault ratio(\n)
64	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f(\n)
128	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f(\n)
256	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f(\n)
512	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f(\n)

Total elapsed time %.4f sec(\n)

# The scenario of program

- Your output should be exactly the same as follows
- Test case : whois www.nctu.edu.tw

```
ubuntu@ubuntu:~/C-project/OS/oshw5$ ./osSummerHW5 ./testcase/whois.txt
LFU policy:
Frame   Hit           Miss           Page fault ratio
64      689115         83143         0.1076622062
128     729104         43154         0.0558802887
256     771124         1134          0.0014684212
512     771921         337           0.0004363827
Total elapsed time 0.1785 sec

LRU policy:
Frame   Hit           Miss           Page fault ratio
64      771056         1202          0.0015564746
128     771663         595           0.0007704679
256     771909         349           0.0004519215
512     771921         337           0.0004363827
Total elapsed time 0.1916 sec
```

- Your output should be exactly the same as follows
- Test case : zipfian

```
ubuntu@ubuntu:~/C-project/OS/oshw5$ ./osSummerHW5 ./testcase/zipf.txt
LFU policy:
Frame    Hit           Miss           Page fault ratio
64       7880135         2119865         0.2119865000
128      8382340         1617660         0.1617660000
256      8807903         1192097         0.1192097000
512      9134227         865773          0.0865773000
Total elapsed time 2.5885 sec

LRU policy:
Frame    Hit           Miss           Page fault ratio
64       7215568         2784432         0.2784432000
128      7880805         2119195         0.2119195000
256      8425224         1574776         0.1574776000
512      8871063         1128937         0.1128937000
Total elapsed time 2.0990 sec
```

# Correctness

- The TAs will prepare **another workload** to validate your implementation
- Except the total elapsed time, your results should be exactly the same as ours
- Do not use linear search in anywhere of your program; otherwise, you will receive a score penalty



# More details

- Total request #  $\leq 10$  millions
- Max page # == 0xffff
- The path+file name of the trace file is an argument of your program (see the screen shot), do not hard-coding the pathname of the trace file
- For each iteration, you should open the file, run the simulation, print the result and close the file
- Do not store the trace data to speed up the next iteration
- Use `gettimeofday()` to get the total elapsed time

# Header of your .c or .cpp

```
/*
```

```
Student No.: <your student id>
```

```
Student Name: <your name>
```

```
Email: <your email>
```

```
SE tag: xnxctxuxoxsx
```

```
Statement: I am fully aware that this program is not  
supposed to be posted to a public server, such as a  
public GitHub repository or a public web page.
```

```
*/
```

# About Submission

- We will compile your code on Ubuntu 14.04 , 16.04, 18.04 or CS linux work station
- You should name your program into <student\_id>.c or <student\_id>.cpp
- How we compile your program:  
“ gcc -o <student\_id>.out <student\_id>.c ” or  
“ g++ -o <student\_id>.out <student\_id>.cpp ”
- If you use a make file, please create a directory named as <student\_id>, put all files into the directory. Compress the directory and upload the .zip file to the E3 platform